

Latency Distribution Across Architectures

